

Machine Learning and Deep Learning

Lecture-03

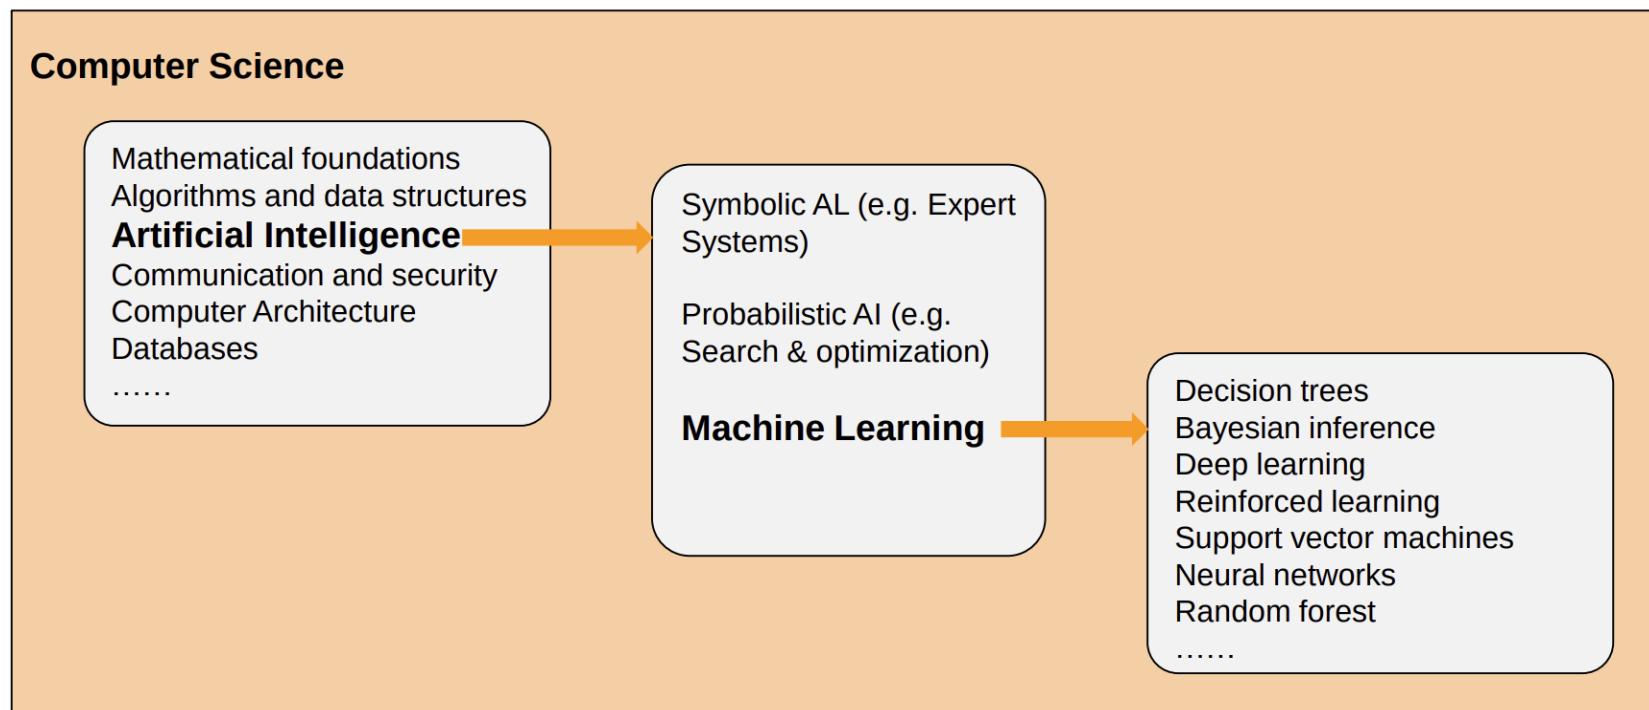
By: Somnath Mazumdar
Assistant Professor
sma.digi@cbs.dk

Outline

- Introduction
- Unsupervised Machine Learning
 - K-Means
 - Density Clustering DBSCAN
 - Gaussian Mixture Models

Fundamentals of Machine Learning

- Artificial Intelligence (AI) is a branch of Computer Science that uses algorithms and techniques to mimic human intelligence.
- Machine Learning (ML) is one of several AI techniques for sophisticated cognitive tasks.



Fundamentals of Machine Learning

Traditional AI techniques



Static, Rule Based, No generalization

Machine Learning



Dynamic, Data driven, Generalization

Fundamentals of Machine Learning

- Three Approaches:
 - Example - Excelling at playing chess game.
- **Symbolic AI**: “Let us sit down with the world’s best chess player, Magnus Carlsen, and put his knowledge into a computer program”.
- **Statistical AI**: “**Let us simulate all the different possible moves** and the associated outcomes at each single step and go with the most likely to win”.
- **Machine Learning Approach**: “**Let us show millions of examples** or real life and simulated games (won and lost) to the program, and **let it learn** from experience”.

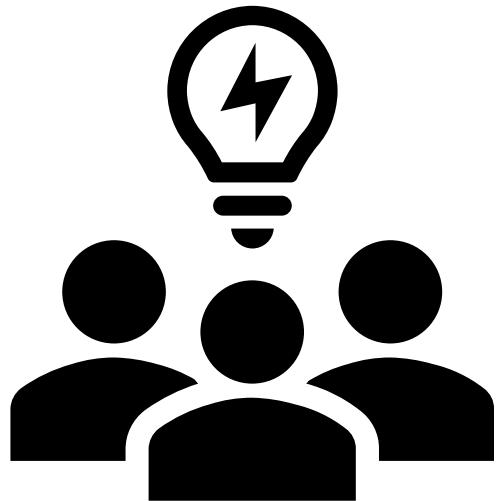


Fundamentals of Machine Learning

- ML is good at solving 2 types of problems where other AI techniques fail:
 1. Tasks programmers can't describe (Handwriting, Cognitive Reasoning).
 2. Complex multidimensional problems that **can't be solved** by numerical reasoning (Weather Forecasting, Health Care Outcomes).
- Unsupervised learning: Want to find unknown structures or trends.
- Supervised learning: Already know answers we want (found in past or completed data).

What do you think?

- Consider a financial startup, 'X' have a large data set, but it does not have any data scientists or ML-guys but has computer science guys. How should they proceed? They do not know what questions to ask!!
- Condition: They will wait to hire a new data scientist!!
- Hints: Supervised vs Unsupervised model – which one they should apply and why?



2 minutes to think

What do you think?

- Unsupervised model
 1. Examine various properties of data.
 2. No data labeling.
 3. Simpler model.
 4. Cheaper to run.

Unsupervised Learning

- Unsupervised Learning: Search for structure in data. Unknown targets.
 - User inputs data with undefined answers.
 - Machine finds useful information hidden in data.
- Example:
 - Cluster analysis: Group into sets: K-Means Clustering, Hierarchical Clustering.
 - Density reduction: Select relevant variables.
 - Dimension Reduction: Principal component analysis.

Three C's of ML

- Three C's of ML:
 1. Collaborative filtering
 2. Clustering
 3. Classification
- **Collaborative filtering** is a technique for recommendations.
 - It's one primary type of recommender system.
 - Can use the same algorithm to recommend practically anything – Movies (Netflix).
- Amazon uses it to recommend a variety of products.

Customers who viewed this item also viewed

Page 1 of 6

The screenshot shows a row of seven recommended products under the heading "Customers who viewed this item also viewed". Each item includes a small image, the product name, a short description, its rating, the number of reviews, and the price. Navigation arrows are at the top and bottom of the list.

Product	Description	Rating	Reviews	Price
Logitech C930 HD Webcam	1080p and USB Port Black/Silver and Zone Wireless Business...	★★★★★	1,134	€195.00
Trust GXT 1160 Vero Full HD 1080p Webcam	Black	★★★★★	9	€49.99
1080P Full HD Webcam, Webcam with Microphone with Data Protection Cover,...	Black	★★★★★	2	€15.99
AUKEY Webcam 1080P Full HD with Stereo Microphone Web Camera for Video Chat and...	Black	★★★★★	4,654	€49.99
Streaming Camera PC 1080P Gaming Live Webcam with Studio Style Ring Light, Dual...	Black	★★★★★	123	€87.99
Blue Microphones	★★★★★	397	€177.88	
Logitech C925e Webcam + Logitech H800 Wireless Headset, Black, Black	Black	★★★★★	296	€167.95

Three C's of ML

- **Clustering** algorithms discover structure in collections of data.
 - Where no formal structure previously existed.
 - They discover what clusters ('groupings'), naturally occur in data
 - By examining various properties of input data.
- Clustering is often used for exploratory analysis: **Divide and Explore !!**
- Applications:
 - Market segmentation → Targeted marketing.
 - Finding related news articles → Google News.

Three C's of ML

- **Classification** is a form of ‘supervised’ learning.
 - Requires training with data that has known labels
 - Learns how to label new records based on that information.
- Applications:
 - Spam filtering: Train using a set of spam and non/spam messages --> System will eventually learn to detect unwanted e-mail
 - Risk Analysis: Train using financial records of customers who do/don’t default --> System will eventually learn to identify risk customers

No 'BEST!!'

- There are many candidate algorithms for a problem.
- **No** overall best algorithm.
- Each algorithm has advantages and limitations.
 - Best approach = simple algorithm + lots of data.
- Algorithm choice is often related to **data** volume <Bigdata>.
- Some scale better than others.
 - Most algorithms offer better results as volume increases.

Applications

- Dimensionality reduction: As a preprocessing step to reduce number of dimensions while maintaining meaningful properties of input data.
- Outlier detection: Identify instances with low affinity to all clusters.
- Semi-supervised learning: When few labels are available, perform clustering and propagate the labels to all the instances in the same cluster.

Supervised Learning

Unsupervised Learning

Continuous *Discrete*

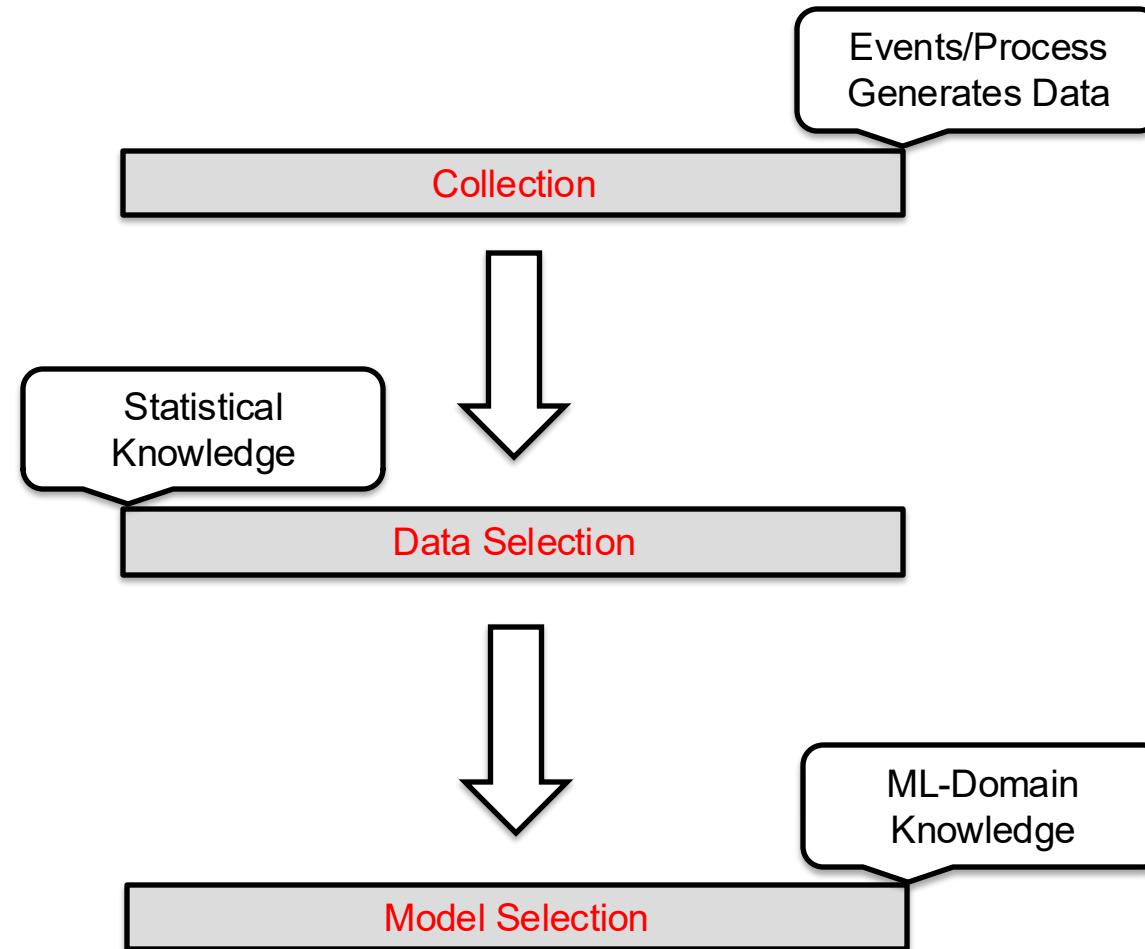
classification or
categorization

clustering

regression

dimensionality
reduction

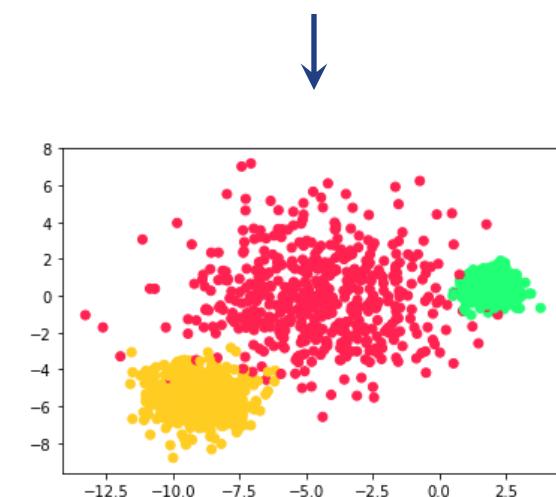
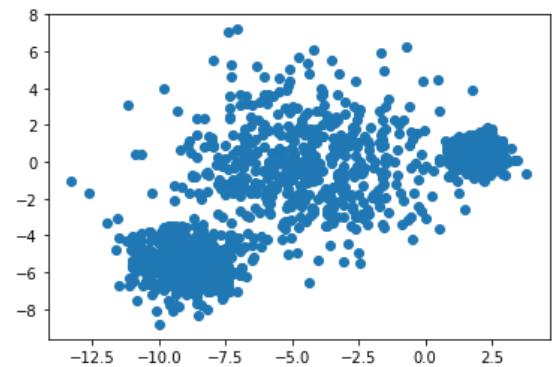
Before you Start..



Clustering

Unsupervised learning

- Labels are not provided, and the **algorithm learns how to group the data without being explicitly told** what the groups are
- E.g., grouping a webstore's customers into segments or finding an outlier in transactions
- Example algorithms:
 - Clustering: K-means, DBSCAN, Hierarchical
 - Dimensionality reduction: Principal components analysis (PCA).



Example

- Often data comes without labels.
- Producing labeled training data can be difficult and time consuming.

Exclusive: OpenAI Used Kenyan Workers on Less Than \$2 Per Hour to Make ChatGPT Less Toxic



<https://time.com/6247678/openai-chatgpt-kenya-workers/>

To build that safety system, OpenAI took a leaf out of the playbook of social media companies like Facebook, who had already shown it was possible to build AIs that could detect toxic language like hate speech to help remove it from their platforms. The premise was simple: feed an AI with labeled examples of violence, hate speech, and sexual abuse, and that tool could learn to detect those forms of toxicity in the wild. That detector would be built into ChatGPT to check whether it was echoing the toxicity of its training data, and filter it out before it ever reached the user. It could also help scrub toxic text from the training datasets of future AI models.

Clustering

- Input: Set of objects described by features x_i .
- Output: An assignment of objects to “groups”.
- Unlike classification, we are not given “groups”
 - Clustering algorithm **must discover "groups"**.
- Examples of groups when clustering the webstore’s customers:
 - Wealthy customers
 - Discount hunting customers

Clustering

- Goals:
 - Objects in the same group should be ‘similar’.
 - Objects in different groups should be ‘different’.
- ‘Best’ clustering is hard to define:
 - No test error.
 - No ‘best’ method.
- Why cluster?
 - To know what the groups are.
 - A ‘prototype’ example for each group.
 - To find the group for a new example x_i .
 - To find objects related to a new example x_i .

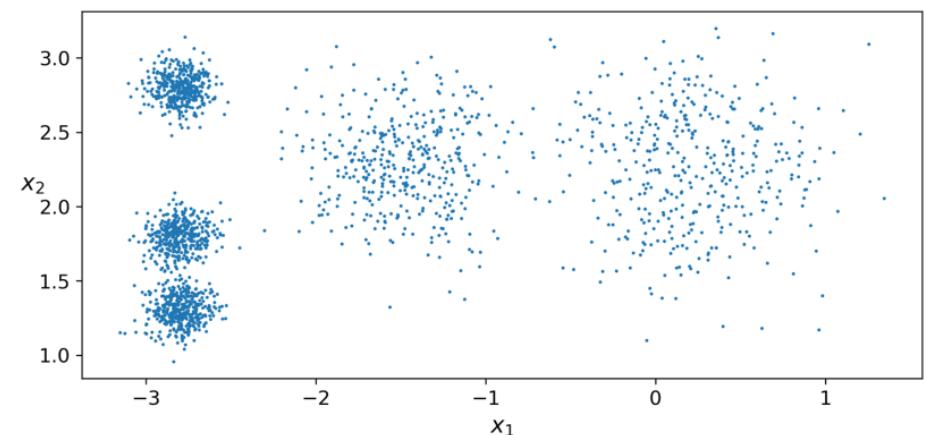
Applications of Clustering

- Clustering is used in a wide variety of applications:
 - For customer segmentation
 - For data analysis
 - As a dimensionality reduction technique
 - For anomaly detection (outlier detection)

K-Means

K-Means

- K-Means algorithm is a simple and efficient algorithm.
- It clusters dataset quickly, often in few iterations.
- Proposed by Stuart Lloyd at Bell Labs, 1957 but published in 1982.
- In 1965, Edward W. Forgy published virtually the same algorithm, it is referred to as Lloyd–Forgy.

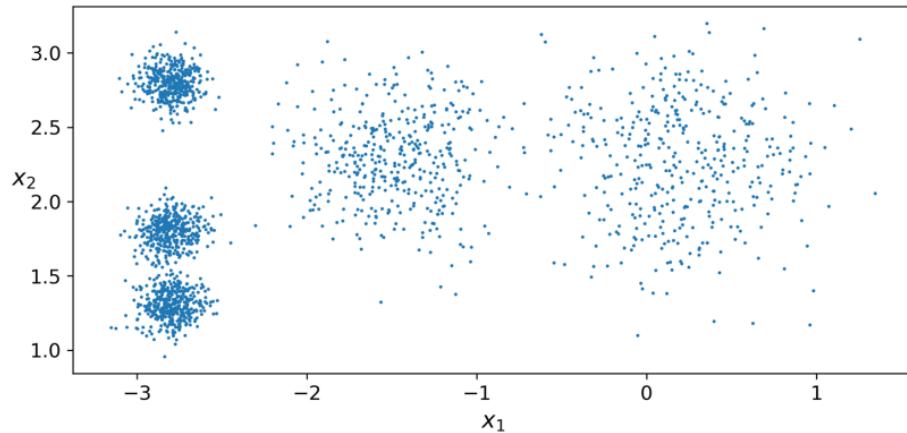


An unlabeled dataset composed of five blobs of instances

K-Means Algorithm

- Suppose centroids are given:
 - Can easily label all instances in dataset by assigning each of them to cluster whose centroid is closest.
- All instance labels were given:
 - Easily locate all centroids by computing mean of instances for each cluster.
- If neither labels nor centroids are given?
 - Start by placing centroids randomly (k -instances).
 - Label instances, update centroids, continue until centroids stop moving.
- It guaranteed to converge in a finite number of steps.

K-Means Algorithm



Unlabeled dataset composed of five blobs of instances



Q: Find each blob's center

```
from sklearn.cluster import KMeans  
k = 5  
kmeans = KMeans(n_clusters=k)  
y_pred = kmeans.fit_predict(X)
```

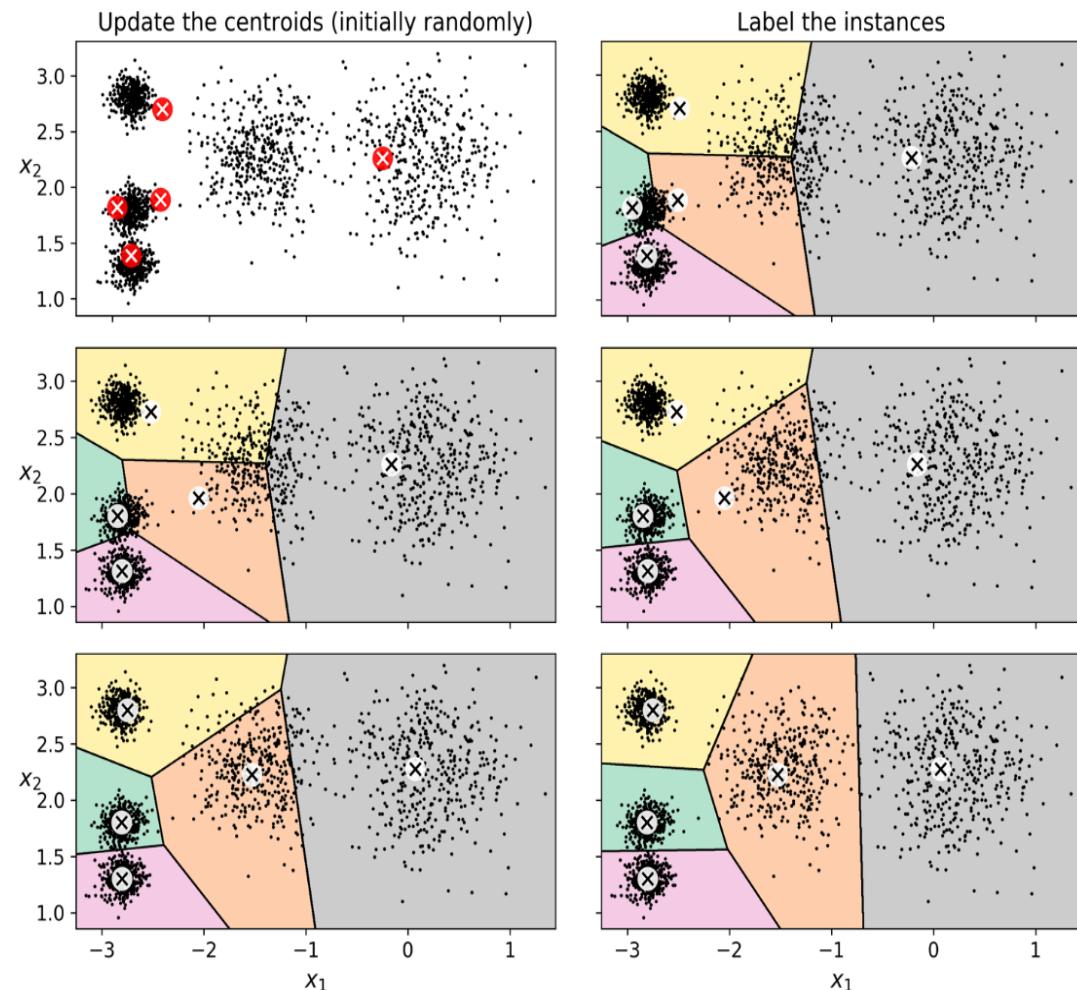
Train a K-Means

```
>>> kmeans.cluster_centers_  
array([[-2.80389616,  1.80117999],  
       [ 0.20876306,  2.25551336],  
       [-2.79290307,  2.79641063],  
       [-1.46679593,  2.28585348],  
       [-2.80037642,  1.30082566]])
```

Five centroids that the algorithm found

K-Means Algorithm

1. Centroids are initialized randomly (top left)
 2. Instances are labeled (top right)
 3. Centroids are updated (center left)
 4. Instances are relabeled (center right)
 5. and so on.
-
- Algorithm has reached an optimal clustering in 3 iterations.

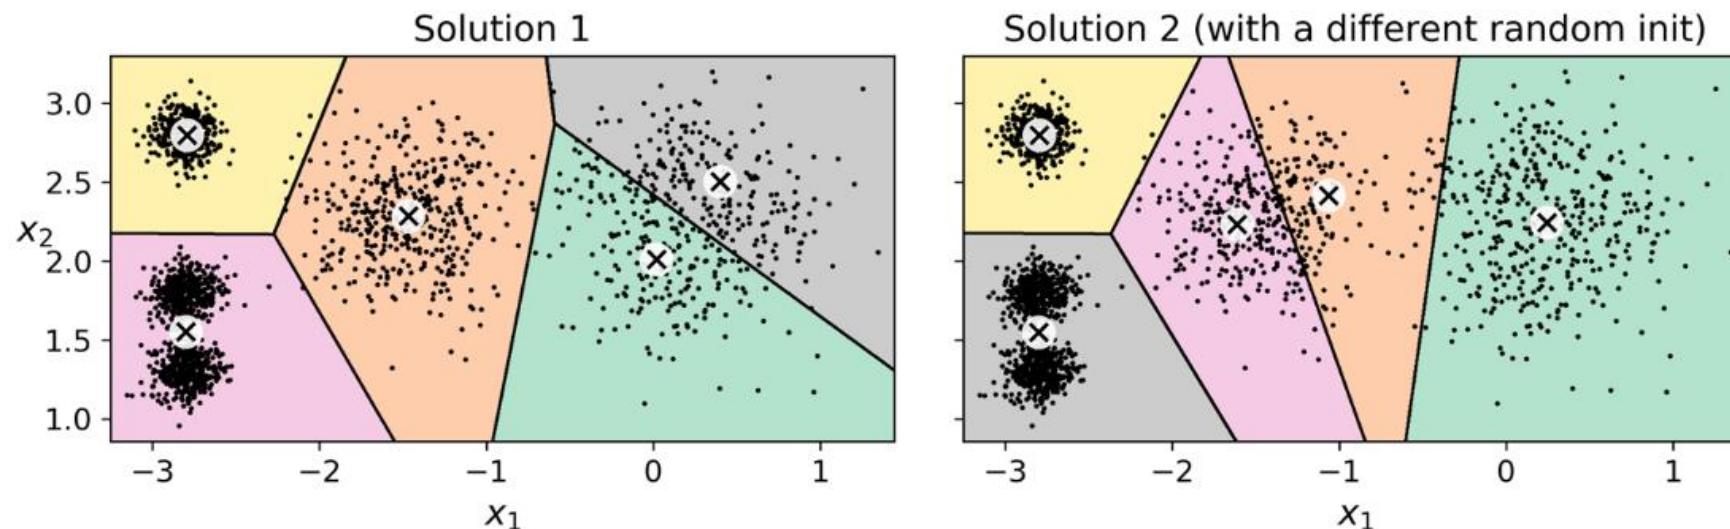


K-Means Algorithm

- Computational complexity is generally linear regarding number of:
 - m (instances), k (number of clusters), and n (number of dimensions).
 - True when data has a clustering structure.
 - If it does not, then in worst-case scenario complexity can increase exponentially with number of instances.
- K-Means is one of the **fastest** clustering algorithms.

K-Means Algorithm

- Algorithm is guaranteed to converge, BUT may **not converge to right solution** (i.e., converge to a local optimum):
 - Depends on centroid initialization.



K-Means Algorithm

- Centroid initialization methods:
 - Set `init` hyperparameter to a NumPy array containing the list of centroids.
 - Set `n_init` to 1: --> if you know approximately where the centroids should be.
 - Run the algorithm multiple times with different random initializations and keep the best solution.

```
good_init = np.array([[-3, 3], [-3, 2], [-3, 1], [-1, 2], [0, 2]])  
kmeans = KMeans(n_clusters=5, init=good_init, n_init=1)
```

- `n_init` is hyperparameter.
- Algorithm will run 10 times

By default, `n_init` hyperparameter set to 10

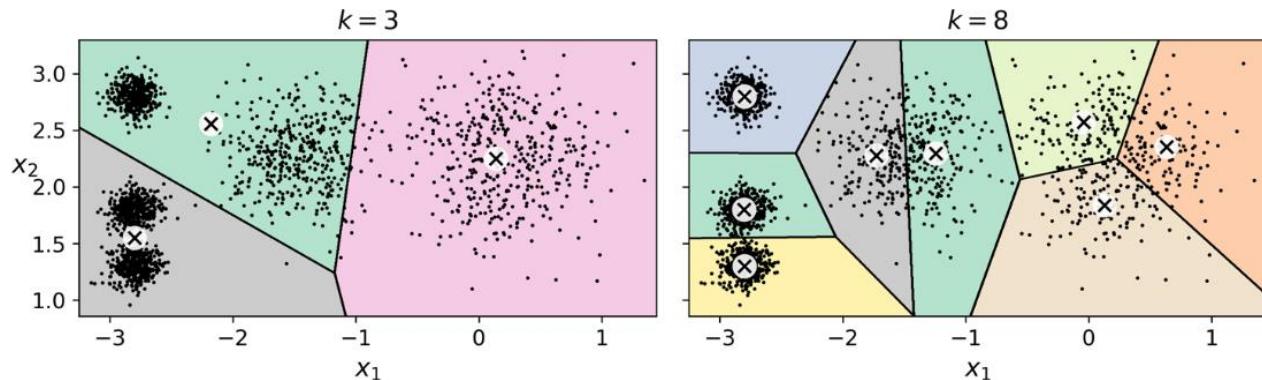
K-Means Algorithm: Mini-Batch

- Mini-batch: Instead of using full dataset at each iteration, use mini-batches, moving the centroids just slightly at each iteration.
 - Speeds up the algorithm typically by a factor of 3x or 4x.
 - Possible to cluster huge datasets that do not fit in memory.

```
from sklearn.cluster import MiniBatchKMeans  
  
minibatch_kmeans = MiniBatchKMeans(n_clusters=5)  
minibatch_kmeans.fit(X)
```

Silhouette Score

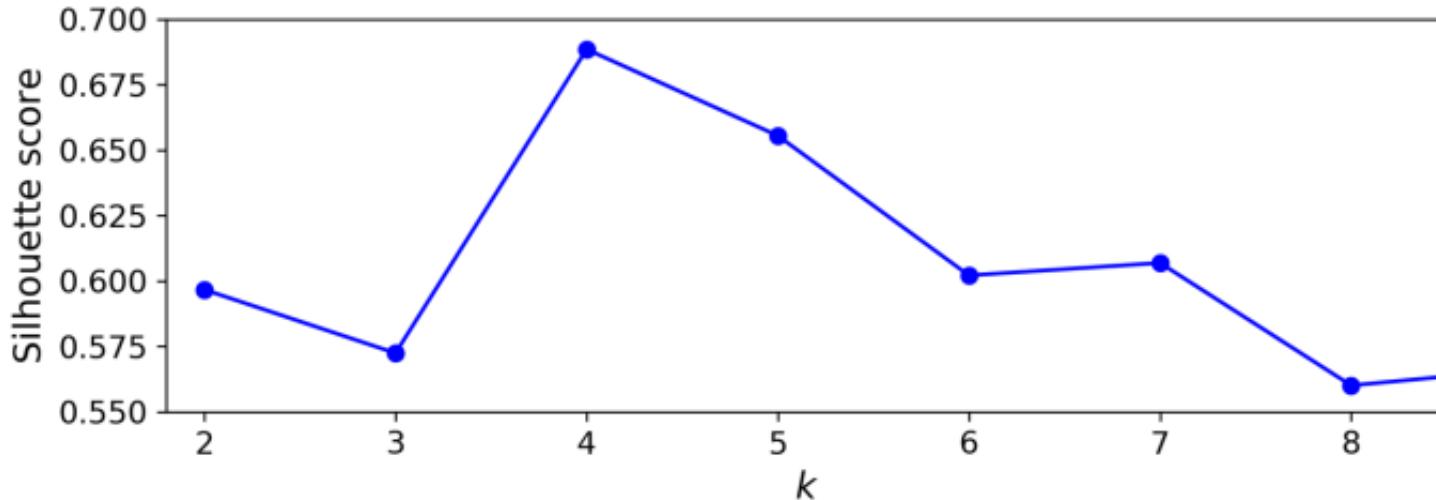
- Silhouette score to find **optimal number of clusters**.
- It is the mean silhouette coefficient over all the instances.
- $(b-a)/\max(a,b)$ where a: mean distance to other instances in same cluster, b: mean nearest-cluster distance.
- Silhouette coefficient vary between -1 and $+1$.



- k is too small, separate clusters get merged (left)
- k is too large, clusters get chopped into multiple pieces (right)

Silhouette Score

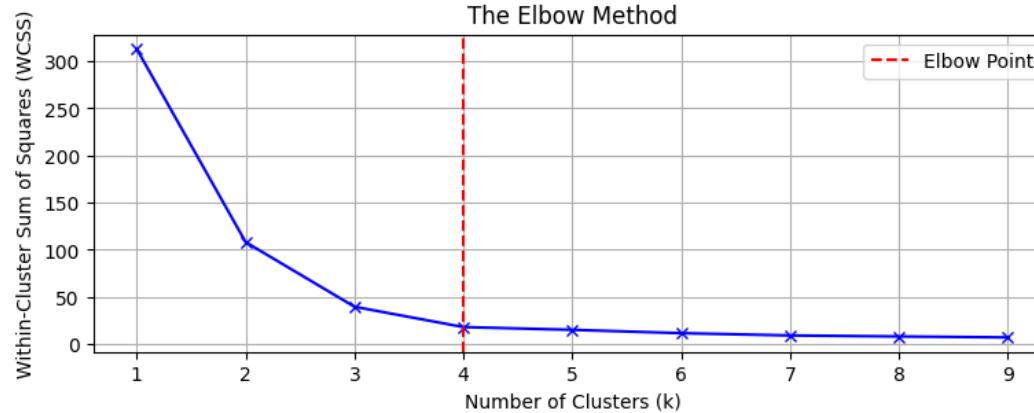
```
from sklearn.metrics import silhouette_score  
silhouette_score(X, kmeans.labels_)
```



$k = 4$ is very good, $k = 5$ is quite good much better than $k = 6$ or 7 .

Elbow Method

- Finds optimal K value.
 - It is the point at which the graph forms an elbow.
- Steps:
 - Iterate over a range (1 to n) of k values, n is a hyper-parameter.
 - For each k, calculate the Within-Cluster Sum of Squares (WCSS).
 - WCSS tells us how well data points are clustered around their respective centroids.



- Prefer the Silhouette score because:
 - Silhouette score is more effective to find number of K when the elbow method does not show the elbow point.
 - Elbow method can produce more ambiguous result.

Issues of K-Means

- Necessary to run algorithm several times to avoid sub-optimal solutions.
- Need to specify the number of clusters.
- Does not behave very well when the clusters have:
 - Varying sizes
 - Different densities
 - Non-spherical shapes.
- Must scale input features before you run K-Means.
- Scaling features improve the performance.

Density Clustering

Density Clustering

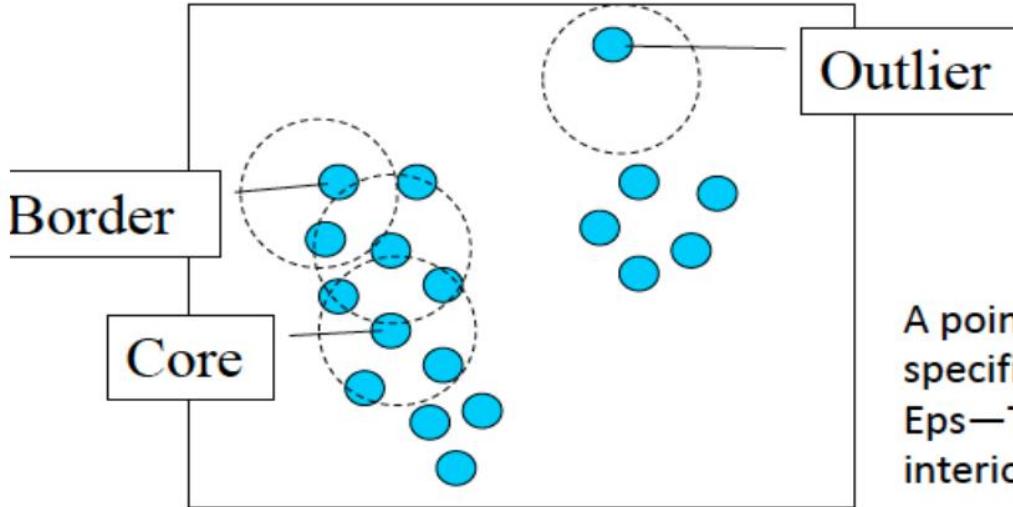
- It is an unsupervised learning method.
- Identify distinctive groups/clusters in input data.
- Idea: A cluster in a data space is a **contiguous region** of high point density, separated from other such clusters by contiguous regions of low point density.
- Data points in separating regions of low point density are typically considered **noise/outliers**.



Density Clustering

- Locates regions of high density that are separated from one another by regions of low density.
- Main principles
 - Maximum radius of the neighbourhood (Eps)
 - Minimum number of points ($MinPts$) in an Eps neighbourhood of a point: $N_{Eps}(p) : \{q \in D \text{ s.t. } dist(p, q) \leq Eps\}$
 - Idea: Density of neighbourhood must exceed some threshold.
 - Shape of a neighbourhood depends on $dist$ function.

Core, Border and Noise/Outlier



Given ϵ and $MinPts$, categorize the objects into three exclusive groups.

A point is a **core point** if it has more than a specified number of points ($MinPts$) within Eps —These are points that are at the interior of a cluster.

$$\epsilon = 1\text{unit}, MinPts = 5$$

A **border point** has fewer than $MinPts$ within Eps , but is in the neighborhood of a core point.

A **noise point** is any point that is not a core point nor a border point.

Density Clustering

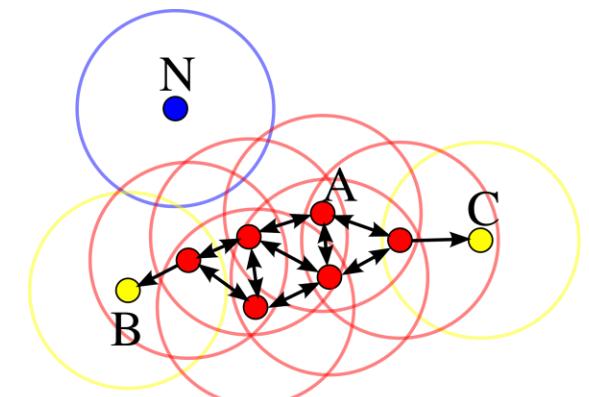
- + Discovers clusters of arbitrary shapes.
- + Handles noise.
- + Needs density parameters as termination condition.

- Cannot handle varying densities.
- Sensitive to parameters.
- Sampling affects density measures.

DBSCAN

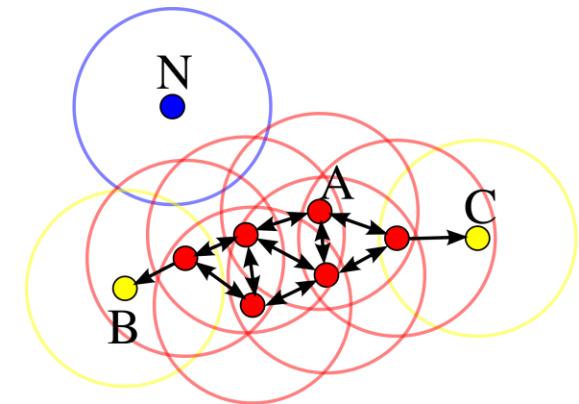
- DBSCAN - Density-Based Spatial Clustering of Applications with Noise.
 - Two parameters Eps and Minpts are set according to experience !!
 - Finds core samples of high density and expands clusters from them.
 - Good for data which contains clusters of similar density.
 - Worst case memory complexity is $O(n^2)$ [when eps is large and minPts is small].

minPts = 4. Red points are core points, because the area surrounding these points in an ϵ radius contain at least 4 points (including the point itself). Yellow are not core points. Point N is a noise point.



DBSCAN Algorithm

- Defines clusters as continuous regions of high density.
- How DBSCAN works?
 1. For each instance, algorithm counts how many instances are located within a small distance ϵ (epsilon) from it. Region is called instance's ϵ -neighborhood.
 2. If an instance has at least `min_samples` instances in its ϵ -neighborhood, then it is considered a core instance.
 3. All instances in the neighborhood of a core instance belong to same cluster. Neighborhood may include other core instances.
 4. Any instance that is not a core instance and does not have one in its neighborhood is considered an anomaly.

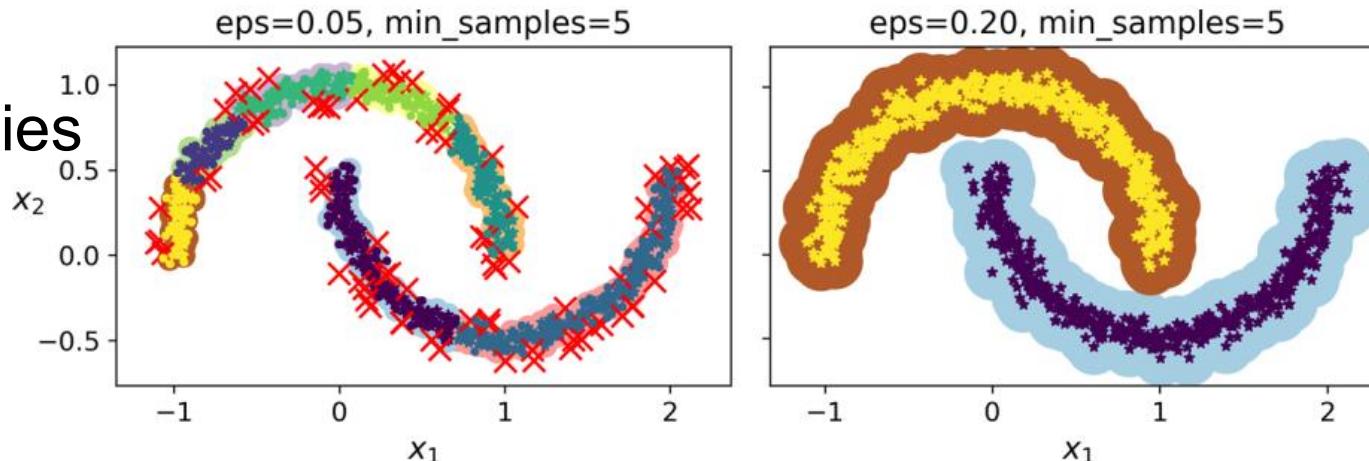


It works well if all the clusters
are dense enough!!

Example

```
from sklearn.cluster import DBSCAN  
from sklearn.datasets import make_moons<--moons dataset  
  
X, y = make_moons(n_samples=1000, noise=0.05)  
dbscan = DBSCAN(eps=0.05, min_samples=5)  
dbscan.fit(X)
```

- Showing anomalies & 7 different clusters.

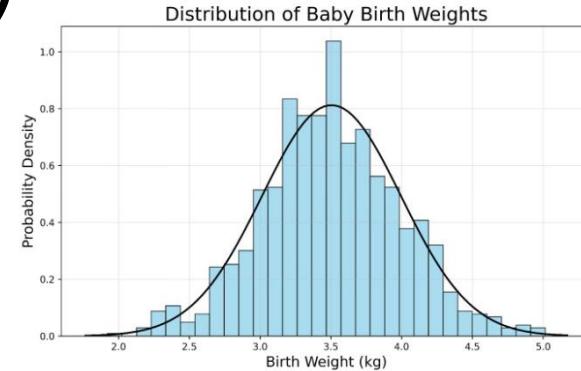


- Widen each instance's neighborhood by increasing eps to 0.2

Gaussian Mixture Models

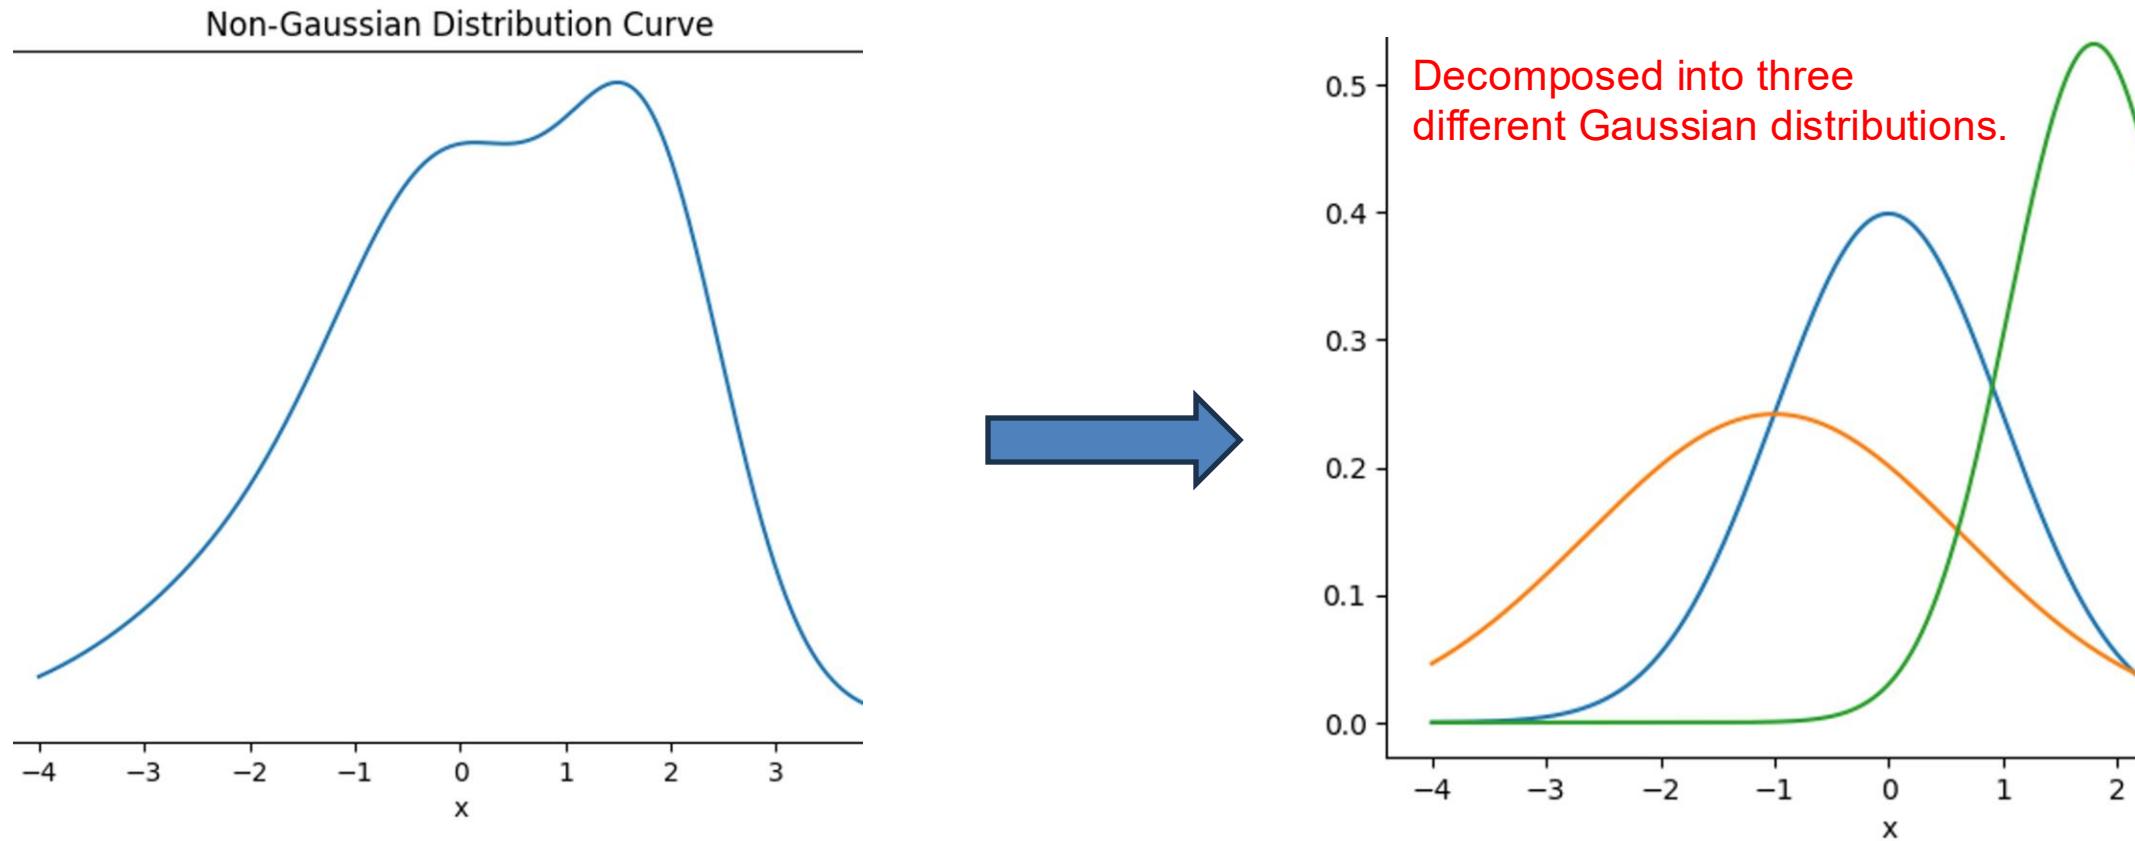
Gaussian Mixture Model (GMM)

- A single Gaussian distribution ==> “normal distribution”
- It describes many kinds of natural phenomena.
 - Example: Distribution of birth weights for full-term babies in a large population.
- Single Gaussian distribution isn’t suitable for modeling datasets with multiple clusters of data or those with a significant skew or heavy tails.
- GMM is a probabilistic model that represents data as a combination of several Gaussian distributions, each with its own mean and variance, weighted by a mixing coefficient.
- GMM is used for clustering and density estimation, since they can capture complex, multimodal distributions.
- GMMs are powerful but they rely on Gaussian assumptions.
- For binary or categorical data, GMMs do not work well



Gaussian Mixture Model (GMM)

- GMMs have many real-world applications beyond clustering:
 - E.g., segmentation, density estimation, anomaly detection and pattern recognition.



GMM works for higher dimensional distributions as well.

Gaussian Mixture Model (GMM)

- When used as a clustering algorithm, each Gaussian in the mixture model has three key parameters:
 - Mean vector: center of cluster. In a 1D distribution this will be a single valued vector. (In an n-D distribution --> n-valued vector).
 - Covariance matrix: Shape of Gaussian itself. In a 1D distribution this will be a single value, (In an n-D distribution --> nxn matrix).
 - Mixing weight: Probability that a randomly chosen data point was generated by component.
- Goal: To estimate both parameters of each Gaussian distribution in the model and which of those Gaussians each data point belongs to.
- This variable is “latent” because it’s a hidden (or unobserved) variable that can be learned from the model.
- Each Gaussian component is weighted by a mixing coefficient, which represents an estimation of how much each distribution affects the location of that specific datapoint.

Gaussian Mixture Model (GMM)

- Feature engineering: while some ML algorithms (XGBoost) can enable a model to learn a variety of input feature distributions, others are stricter in their requirements.
 - Linear and logistic regression typically expect features to be normally distributed (and may not function well if the data is multimodal).
- Unsupervised classification: GMM works similarly to the K-means algorithm but allows probabilistic determination of class belonging, unlike K-means, where the output is a binary metric. This can be especially beneficial to use cases that require custom thresholds for categorization or that require a probabilistic output.
- Anomaly detection: A multivariate Gaussian distribution can be used to identify data points that have low probability of following one or more Gaussian distributions.

Gaussian Mixture Model (GMM)

- When used for clustering, GMMs are similar to k-means clustering but have several key differences. GMMs give probabilities of belonging to each cluster.
- Since clusters can be both elliptical and overlapping, GMMs are often more flexible and allow for more uncertainty in cluster boundaries.
- An extension of GMM is the variational autoencoder (VAE), which is a generative model that learns flexible latent distributions.
- A VAE uses a probabilistic encoder-decoder framework to learn latent representations in the same way that a GMM assigns mixture weights for each data point.

Conclusion

- Key Takeaways
 - Unsupervised learning discovers structure in data without labeled outputs.
 - Clustering is a foundational tool for exploratory analysis and segmentation.
 - There is no single “best” algorithm (method choice depends on data characteristics)
 - Simpler models combined with sufficient data often perform well.
 - Understanding assumptions is critical to correct algorithm usage.
 - Use unsupervised learning when Labels are unavailable or costly.
- K-Means
 - Fast and scalable
 - Requires predefined number of clusters
 - Sensitive to initialization and feature scaling
 - Best for spherical, similarly sized clusters

Conclusion

- DBSCAN
 - Identifies clusters of arbitrary shape.
 - Handles noise and outliers naturally.
 - Sensitive to density parameters.
 - Struggles with varying cluster densities.
- Gaussian Mixture Models (GMM)
 - Probabilistic cluster assignment.
 - Handles overlapping and elliptical clusters.
 - Assumes Gaussian data distributions.
 - More flexible but computationally heavier.
- Always:
 - Inspect data distributions
 - Scale features where required
 - Validate clustering quality (e.g., Silhouette score)